

Bayesian modelling of football outcomes: using the Skellam's distribution for the goal difference

DIMITRIS KARLIS[†] AND IOANNIS NTZOUFRAS[‡]

*Department of Statistics, Athens University of Economics and Business,
76 Patission Street, 10434 Athens, Greece*

[Received on 26 July 2007; accepted on 14 August 2008]

Modelling football match outcomes is becoming increasingly popular nowadays for both team managers and betting fans. Most of the existing literature deals with modelling the number of goals scored by each team. In this paper, we work in a different direction. Instead of modelling the number of goals directly, we focus on the difference of the number of goals, i.e. the margin of victory. Modelling the differences instead of the scores themselves has some major advantages. Firstly, we eliminate correlation imposed by the fact that the two opponent teams compete each other, and secondly, we do not assume that the scored goals by each team are marginally Poisson distributed. Application of the Bayesian methodology for the Skellam's distribution using covariates is discussed. Illustrations using real data from the English Premiership for the season 2006–2007 are provided. The advantages of the proposed approach are also discussed.

Keywords: goal difference; overdispersion; Poisson difference; Skellam's distribution; soccer; zero-inflated models.

1. Introduction

In recent years, an increasing interest has been observed concerning models related to football (soccer). The increasing popularity of football modelling and prediction is mainly due to two distinct reasons. Firstly, the market related to football has increased considerably the past years; modern football teams are profitable companies usually with large investments and budgets, while the interest of the sports fans concerning football is extremely large. The second reason is betting. The amount spent on bets has increased dramatically in Europe. As a result, the demand for models which provide good predictions for the outcome of a football game arises. Since bets are becoming more complicated, more complicated and more refined models are needed.

The Poisson distribution has been widely used as a simple modelling approach for describing the number of goals in football (see, e.g. [Lee, 1997](#)). This assumption can be questionable in certain leagues where overdispersion (sample variance exceeds the sample mean) has been observed in the number of goals. In addition, empirical evidence has shown a (relatively low) correlation between the goals in a football game. This correlation can be incorporated in the model (see [Maher, 1982](#); [Dixon & Coles, 1997](#); [Karlis & Ntzoufras, 2003](#)).

This paper moves in a different direction. Instead of modelling jointly the number of goals scored by each team, we focus on the goal difference. By this way, we remove the effect of the correlation between

[†]Corresponding author. Email: karlis@aub.gr

[‡]Email: ntzoufras@aub.gr

the scoring performance of the two competing teams, while the proposed model does not assume Poisson marginals. The model can be used to predict the outcome of the game as well as for betting purposes related to the Asian handicap. However, the model cannot predict the final score.

We proceed using the Bayesian approach concerning the estimation of the model parameters along the lines of [Karlis & Ntzoufras \(2006\)](#). The Bayesian approach is suitable for modelling sports outcomes in general, since it allows the user to incorporate any available information about each game via the prior distribution. Information that can be incorporated in the model can be based on historical knowledge or data, weather conditions or the fitness of a team. Finally, the Bayesian approach naturally allows for predictions via the predictive distribution. This allows to predict future games and produce a posterior predictive distribution for future scores and outcomes or even reproduce the whole tournament and produce quantitative measures concerning the performance of each team.

The remaining of the paper proceeds as follows: Section 2 describes the proposed model and provides some properties and interesting points that will assist us to understand the behaviour of the model and its interpretation. A zero-inflated model is also described in order to capture the (possible) excess of draws in a league. In Section 3, we describe the Bayesian approach used to estimate the model. A real-data application using the results of the English Premier League for the season 2006–2007 is described in Section 4. Finally, concluding remarks and further work are discussed in Section 5.

2. The model

2.1 Derivation

Consider two discrete random variables X and Y and their difference $Z = X - Y$. The probability function of the difference Z is a discrete distribution defined on the set of integer numbers $\mathcal{Z} = \{\dots, -2, -1, 0, 1, 2, \dots\}$. Although publications concerning distributions defined on \mathcal{Z} are rare, the difference of two independent Poisson random variables has been discussed by [Irwin \(1937\)](#) for the case of equal means and [Skellam \(1946\)](#) for the case of different means.

The Skellam's distribution (or Poisson difference (PD) distribution) is defined as the distribution of a random variable Z with probability function

$$f_{\text{PD}}(z|\lambda_1, \lambda_2) = P(Z = z|\lambda_1, \lambda_2) = e^{-(\lambda_1 + \lambda_2)} \left(\frac{\lambda_1}{\lambda_2}\right)^{z/2} I_{|z|}(2\sqrt{\lambda_1 \lambda_2}), \quad (2.1)$$

for all $z \in \mathcal{Z}$, $\lambda_1, \lambda_2 > 0$, where $I_r(x)$ is the modified Bessel function of order r (see Abramowitz & Stegun, 1974, p. 375) given by

$$I_r(x) = \left(\frac{x}{2}\right)^r \sum_{k=0}^{\infty} \frac{\left(\frac{x^2}{4}\right)^k}{k! \Gamma(r + k + 1)}.$$

We will denote as $\text{PD}(\lambda_1, \lambda_2)$ the distribution with probability function given in (2.1). Although the Skellam's distribution was originally derived as the difference of two independent Poisson random variables, it can also be derived as the difference of distributions which have a specific trivariate latent variable structure.

LEMMA For any pair of variables (X, Y) that can be written as $X = W_1 + W_3$ and $Y = W_2 + W_3$ with $W_1 \sim \text{Poisson}(\lambda_1)$, $W_2 \sim \text{Poisson}(\lambda_2)$ and W_3 following any distribution with parameter vector θ_3 , then $Z = X - Y$ follows a $\text{PD}(\lambda_1, \lambda_2)$ distribution.

The proof of the above lemma is straightforward. It is, however, interesting that the joint distribution of X and Y is a bivariate distribution with correlation induced by the common stochastic component in both variables, W_3 . For example, if W_3 follows a Poisson distribution, then the joint distribution is the bivariate Poisson (see [Karlis & Ntzoufras, 2003](#)). In addition, the marginal distributions for X and Y will be Poisson distributed only in the case where W_3 is degenerate at zero or a Poisson-distributed random variable. Therefore, in the general formulation, the marginal distributions of X and Y are not any more Poisson but they are defined as the convolution of a Poisson random variable with another discrete random variable of any distributional form. Thus, the marginal distribution can be overdispersed or even underdispersed relative to a Poisson distribution and, hence, a large portion of the distributional assumptions concerning the number of goals scored by each team is removed. This underlines the efficiency and the flexibility of our proposed model.

Although the above lemma implies that the type of the goal difference distribution will be the same regardless of the existence or the type of association between the two variables, this does not imply that the parameter estimates and their interpretation will be the same. Finally, the trivariate reduction scheme used to define the PD distribution provides a suitable data augmentation scheme that can be used efficiently for constructing the estimation algorithm (see Section 2.2 for details).

The expected value of the $PD(\lambda_1, \lambda_2)$ distribution is given by $E(Z) = \lambda_1 - \lambda_2$, while the variance is $Var(Z) = \lambda_1 + \lambda_2$. Note that, for the range of mean values observed in football games, the distribution cannot be sufficiently approximated by the normal distribution and, hence, inference based on simple normal regression can be misleading. Additional properties of the distribution are described in [Karlis & Ntzoufras \(2006\)](#).

2.2 A model for the goal difference

We can use the PD distribution to model goal difference by specifying as the response variable the goal difference in game i . Hence, we specify

$$Z_i = X_i - Y_i \sim PD(\lambda_{1i}, \lambda_{2i})$$

for $i = 1, 2, \dots, n$, where n is the number of games and X_i and Y_i are the number of goals scored by the home and away team, respectively, in i game. Concerning the model parameters λ_{1i} and λ_{2i} , we adopt the same structure as in simple or bivariate Poisson models used for the number of goals scored by each team (see [Lee, 1997](#); [Karlis & Ntzoufras, 2003](#), respectively). Therefore, we set

$$\log(\lambda_{1i}) = \mu + H + A_{HT_i} + D_{AT_i}, \quad (2.2)$$

$$\log(\lambda_{2i}) = \mu + A_{AT_i} + D_{HT_i}, \quad (2.3)$$

where μ is a constant parameter, H is the home effect, A_k and D_k are the ‘net’ attacking and defensive parameters of team k after removing correlation and HT_i and AT_i are the home and away team competing each other in game i .

Note that for parameters A_k and D_k , we propose to use the sum to zero constraints in order to make the model identifiable. Therefore, we need to impose the constraints

$$\sum_{k=1}^K A_k = 0 \quad \text{and} \quad \sum_{k=1}^K D_k = 0, \quad (2.4)$$

where K is the number of different teams competing each other in the available data set. Under the above parametrization, all parameters have a straightforward interpretation since H is the home effect in a game where two opponent teams have the same attacking and defensive skills, μ is a constant parameter corresponding to the PD parameter for the away team in the same case, while A_k and D_k can be interpreted as deviations of the net attacking and defensive abilities from a team of moderate performance.

2.3 Zero-inflated version of the model

As we have already mentioned, the number of draws (and the corresponding probability) are underestimated by Poisson-based models used for the number of goals scored by each team. For this reason, we can specify the zero-inflated version of Skellam's distribution to model the possible excess of draws. Hence, we can define the zero-inflated Poisson difference (ZPD) distribution as the one with probability function

$$f_{\text{ZPD}}(0|p, \lambda_1, \lambda_2) = p + (1 - p)f_{\text{PD}}(0|\lambda_1, \lambda_2), \quad (2.5)$$

$$f_{\text{ZPD}}(z|p, \lambda_1, \lambda_2) = (1 - p)f_{\text{PD}}(z|\lambda_1, \lambda_2), \quad \text{for } z \in \mathcal{Z} \setminus \{0\}, \quad (2.6)$$

where $p \in (0, 1)$ and $f_{\text{PD}}(z|\lambda_1, \lambda_2)$ is given by (2.1).

3. Bayesian inference

3.1 The prior distributions

To fully specify a Bayesian model, we need to specify the prior distribution. When no information is available, we propose to use normal prior distributions for the parameters of the PD model with mean equal to zero and large variance (e.g. 10^4) to express prior ignorance. For the mixing proportion p used in the zero-inflated version of the proposed model, we propose a uniform distribution defined in the $(0, 1)$ interval.

Nevertheless, the Bayesian approach offers the ability to incorporate external information to our inference via our prior distribution. Also, prior elicitation techniques can be employed in order to produce a prior distribution by extracting information by experts on the topic such as sports analysts and bookmakers. In this case, more general prior distributions can be used. For example, we can use normal prior distributions with small variance centred at a certain value for the parameters of the PD model and a beta prior for the mixing proportion for the zero-inflated model.

Finally, the Bayesian approach can be used sequentially by using the previous fixture posterior distribution as a prior distribution. This would help to update much faster our model.

3.2 The posterior distributions

In the Bayesian approach, the inference is based on the posterior distribution of the model parameters θ . In the PD model, we consider the parameter vector $\theta = (\mu, H, A_2, \dots, A_K, B_2, \dots, B_K)$ and we need to calculate the posterior distribution

$$f(\theta|z) = \frac{f_{\text{PD}}(z|\theta)f(\theta)}{\int f_{\text{PD}}(z|\theta)f(\theta)d\theta},$$

where \mathbf{z} is an $n \times 1$ vector with the observed goal differences, $f(\boldsymbol{\theta})$ is the joint prior distribution which is here defined as the product of independent normal distributions and $f_{\text{PD}}(\mathbf{z}|\boldsymbol{\theta})$ is the model likelihood,

$$f_{\text{PD}}(\mathbf{z}|\boldsymbol{\theta}) = \prod_{i=1}^n f_{\text{PD}}(z_i|\lambda_{1i}, \lambda_{2i}),$$

with $f_{\text{PD}}(z|\lambda_1, \lambda_2)$ given by (2.1) and λ_{1i} and λ_{2i} by (2.2) and (2.3), respectively. The attacking and defensive abilities of the omitted team are simply calculated via the constraints (2.4) and, therefore, A_1 and B_1 will be substituted in the likelihood by

$$A_1 = -\sum_{k=2}^K A_k \quad \text{and} \quad B_1 = -\sum_{k=2}^K B_k.$$

Note that the approach is similar for the zero-inflated version but, in this case, we have to additionally estimate the mixing proportion p . Hence, the posterior is given by

$$f(\boldsymbol{\theta}, p|\mathbf{z}) = \frac{\prod_{i=1}^n f_{\text{ZPD}}(z_i|\lambda_{1i}, \lambda_{2i}, p) f(\boldsymbol{\theta}) f(p)}{\int \int \prod_{i=1}^n f_{\text{ZPD}}(z_i|\lambda_{1i}, \lambda_{2i}, p) f(\boldsymbol{\theta}) f(p) d\boldsymbol{\theta} dp},$$

where $f(p)$ is the prior of the mixing proportion p and $f_{\text{ZPD}}(z|\lambda_1, \lambda_2, p)$ is given by (2.5) and (2.6).

Inference concerning the components of the parameter vector $\boldsymbol{\theta}$ (and p) can be based on the posterior summaries of the marginal posterior distribution (mean, median, standard deviation and quantiles). The above posterior distribution is not analytically tractable. For this reason, we use Markov chain Monte Carlo (MCMC) algorithms to generate values from the posterior distribution and hence estimate the posterior distribution of interest and their corresponding measures of fit.

3.3 The MCMC algorithm

Our approach is based on the sampling augmentation scheme proposed by [Karlis & Ntzoufras \(2006\)](#). Hence, a key element for constructing an MCMC algorithm for the proposed PD and ZPD models is to generate the w_{1i} and w_{2i} augmented data for PD model and additionally the latent binary indicators δ_i for the ZPD model. The first set will be used to specify the observed data z_i as a difference of two Poisson-distributed variables, while the latter will be used in the ZPD model to identify from which component we get the observed difference z_i (i.e. from the PD component or from the inflated one).

Hence, in each iteration of the MCMC algorithm, we

- Generate latent data w_{1i} and w_{2i} from

$$f(w_{1i}, w_{2i}|z_i = w_{1i} - w_{2i}, \lambda_{1i}, \lambda_{2i}) \propto \frac{\lambda_{1i}^{w_{1i}}}{w_{1i}!} \frac{\lambda_{2i}^{w_{2i}}}{w_{2i}!} I(z_i = w_{1i} - w_{2i}),$$

where $I(E) = 1$ if E is true and 0 otherwise.

- Generate latent binary indicators δ_i from

$$f(\delta_i|z_i, \lambda_{1i}, \lambda_{2i}) \cong \text{Bernoulli}(\tilde{p}_i) \quad \text{with} \quad \tilde{p}_i = \frac{p}{p + (1-p)f_{\text{PD}}(z_i|\lambda_{1i}, \lambda_{2i})}.$$

Concerning the simulation of the augmented data (w_{1i}, w_{2i}) used in the PD, we propose to use the following Metropolis–Hastings step:

- If $z_i < 0$ and (w_{1i}, w_{2i}) the current values of the augmented data, then
 - Propose $w'_{1i} \sim \text{Poisson}(\lambda_{1i})$ and $w'_{2i} = w'_{1i} - z_i$.
 - Accept the proposed move with probability $\alpha = \min \left\{ 1, \lambda_{2i}^{(w'_{1i} - w_{1i})} \frac{(w_{1i} - z_i)!}{(w'_{1i} - z_i)!} \right\}$.
- If $z_i \geq 0$ and (w_{1i}, w_{2i}) the current values of the augmented data, then
 - Propose $w'_{2i} \sim \text{Poisson}(\lambda_{2i})$ and $w'_{1i} = w'_{2i} + z_i$.
 - Accept the proposed move with probability $\alpha = \min \left\{ 1, \lambda_{1i}^{(w'_{2i} - w_{2i})} \frac{(w_{2i} + z_i)!}{(w'_{2i} + z_i)!} \right\}$.

Given the augmented data $(w_{1i}, w_{2i}, \delta_i)$, the parameters θ can be generated as in simple Poisson log-models with data $\mathbf{y} = (\mathbf{w}_1^{\text{PD}}, \mathbf{w}_2^{\text{PD}})$, where \mathbf{w}_1^{PD} and \mathbf{w}_2^{PD} are vectors with elements w_{1i} and w_{2i} for which $\delta_i = 0$. The conditional posterior distributions will be given by

$$f(\theta|p, \delta, \mathbf{w}_1, \mathbf{w}_2) \propto \prod_{i=1}^n [f_P(w_{1i}|\lambda_{1i}) f_P(w_{2i}|\lambda_{2i})]^{1-\delta_i} f(\theta)$$

and

$$f(p|\theta, \delta, \mathbf{w}_1, \mathbf{w}_2) \propto p^{\sum_{i=1}^n \delta_i} (1-p)^{n-\sum_{i=1}^n \delta_i} f(p).$$

Note that the PD model is similar to setting all $\delta_i = 0$ for all observations (and $p = 0$, respectively). In the case that we use a beta prior distribution with parameters a and b for p , then the above conditional posterior will be also beta with parameters $\sum_{i=1}^n \delta_i + a$ and $n - \sum_{i=1}^n \delta_i + b$. When we wish to impose additional covariates on the mixing proportion, then the parameters can be generated as in the case of a simple logistic regression model having as a response the latent binary indicators δ .

The above algorithm can be implemented in any programming language or more statistical friendly programming software (such as R and Matlab). Alternatively, we can directly use WinBUGS (Spiegelhalter *et al.*, 2003), a statistical tool for the implementation of Bayesian models using MCMC methodology. Results presented in this article have been reproduced using both R and WinBUGS. The latter is available from the authors upon request.

3.4 Simulating future games and leagues from the predictive distribution

An important feature of Bayesian inference is the predictive distribution. Consider a future game between the home team h and away team a . Hence, we wish to predict a future goal difference $z_{(h,a)}^{\text{pred}}$. This can be done directly using the posterior predictive distribution

$$f(z_{(h,a)}^{\text{pred}}|z) = \int f(z_{(h,a)}^{\text{pred}}|\theta) f(\theta|z) d\theta. \quad (3.1)$$

Note that in the ZPD model, θ is replaced by $\theta' = (\theta, p)$. Moreover, $f(z_{(h,a)}^{\text{pred}}|\theta)$ depends only on parameters μ , H , (A_h, D_h) and (A_a, D_a) (and p in the ZPD model) which are related to the teams competing each other in the game we wish to predict. When we wish to predict goal differences z^{pred}

for $n^{\text{pred}} > 1$ games in which the home team $\text{HT}_k^{\text{pred}}$ competes with the away team $\text{AT}_k^{\text{pred}}$ in k th game (for $k = 1, \dots, n^{\text{pred}}$), then the resulting posterior predictive distribution is given by

$$\begin{aligned} f(z^{\text{pred}} | \text{HT}^{\text{pred}}, \text{AT}^{\text{pred}}, z) &= \int f(z^{\text{pred}} | \text{HT}^{\text{pred}}, \text{AT}^{\text{pred}}, \theta) f(\theta | z) d\theta \\ &= \int \prod_{k=1}^{n^{\text{pred}}} f(z_k^{\text{pred}} | \mu, H, A_{\text{HT}_k^{\text{pred}}}, D_{\text{HT}_k^{\text{pred}}}, A_{\text{AT}_k^{\text{pred}}}, D_{\text{AT}_k^{\text{pred}}}) \\ &\quad \times f(\mu, H, A_2, \dots, A_K, D_2, \dots, D_K | z) d\theta, \end{aligned}$$

where HT^{pred} and AT^{pred} are the vectors of length n^{pred} with the competing teams in the future games we wish to predict.

When using an MCMC algorithm, it is straightforward to generate values of the $z_{(h,a)}^{\text{pred}}$ from the corresponding predictive distribution (3.1) by simply adding the following steps in the MCMC sampler used in the PD model:

- Calculate $\lambda_1^{\text{pred}} = \mu + A_h + D_a + H$ and $\lambda_2^{\text{pred}} = \mu + A_a + D_h$.
- Generate w_1^{pred} and w_2^{pred} from a Poisson distribution with parameters λ_1^{pred} and λ_2^{pred} , respectively.
- Set $z_{(h,a)}^{\text{pred}} = w_1^{\text{pred}} - w_2^{\text{pred}}$ as the generated value from the predictive distribution of interest.

In the above procedure, parameters $\mu, H, (A_h, D_h)$ and (A_a, D_a) will be equal to their corresponding values generated in each iteration of the MCMC algorithm.

For the ZPD model, we need to add the following steps in our sampling algorithm:

- Generate δ^{pred} from a Bernoulli with probability p .
- If $\delta^{\text{pred}} = 1$, then set $z_{(h,a)}^{\text{pred}} = 0$; otherwise proceed as in the PD model.

When modelling sports outcomes, it is usual practice to reproduce the predictive distribution of the ranking table in a league or a tournament. This procedure is very useful and was initially introduced by Lee (1997) using plug-in maximum likelihood estimates from a simple Poisson model. It can be used for two reasons. Firstly, to probabilistically quantify the final outcome of a league or a tournament. This can be used to assess the goodness of fit of the model and the overall performance of specific teams in the league. For example, if the ranking resulting from the predictive distribution is in general agreement with the observed data, this implies a good fit of our model. On the other hand, specific deviations may indicate that specific teams performed better or worse than expected in specific games and hence ended up in a different ranking. Secondly, it can be used to estimate the ranking distribution if the competition had a different structure. For example, when we use data from knock-out tournaments we are able to estimate the rankings if the teams were competing in full season in round robin (league) system. Of course this procedure, with the current model, assumes that the teams' performance is constant across the whole competition.

4. Application: the English Premier League 2006–2007

The data refer to the English Premiership for 2006–2007 season. Data were downloaded from the web page <http://soccornet-akamai.espn.go.com>. For all the parameters of the PD component, we have used

normal prior distributions with zero mean and low precision equal to 10^{-4} (i.e. large variance equal to 10^4) to express our prior ignorance. A uniform prior distribution was used for the mixing proportion p of the ZPD model. All results were produced using 10,000 iterations after discarding additional 1000 iterations as a burn-in period.

A plot of the 95% posterior intervals for all net attacking and defensive parameters for all teams is provided in Figs 1 and 2, respectively. According to this plot, Manchester United, Liverpool, Arsenal, Blackburn and Chelsea had the highest net attacking parameters. Concerning the net defensive parameters, Chelsea had a considerably higher parameter than the rest of the teams followed by Manchester United and Arsenal. If we look at the actual goals scored and conceded by each team, we see that Chelsea and Arsenal had scored more goals than Liverpool. Nevertheless, Liverpool ‘attacking’ parameter is much higher than the corresponding parameters of the other teams. This is due to the properties




FIG. 1. 95% posterior intervals for net attacking coefficients (A_i).

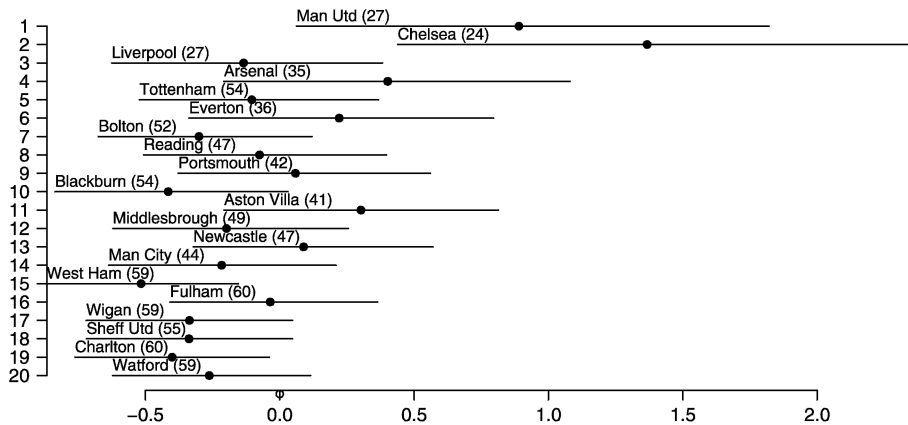


FIG. 2. 95% posterior intervals for net defensive coefficients $[(-1) \times D_i]$; within brackets the observed scored goals.

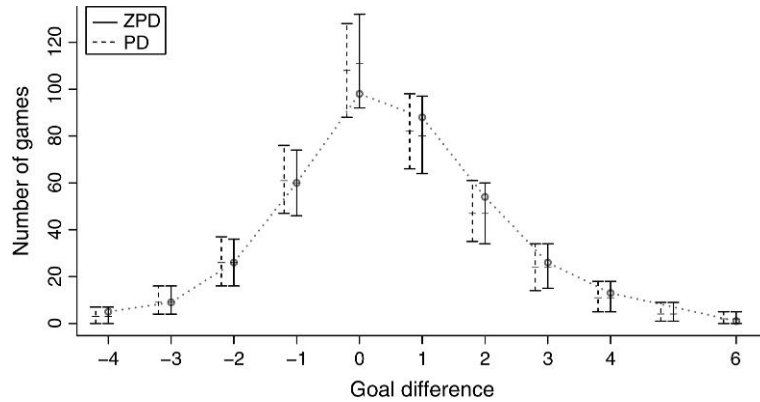


FIG. 3. 95% posterior intervals for predicted differences (dotted line and points indicate the observed differences; ‘-’ indicates the posterior median). No major differences between the two models are observed.

of the PD distribution. Since focus is given in the differences, net attacking and defensive parameters must be interpreted altogether.

Figure 3 depicts the observed and predicted relative frequencies for the goal differences. It is evident that there is a close agreement between the observed and the predicted values. Concerning the number of draws, the PD models slightly overestimate them which comes in contradiction with the results observed when using simple or bivariate Poisson models for the actual number of goals. However, note that the marginal distributions for our model are not necessarily Poisson and this can explain why our model fits better the draws. Nevertheless, the 95% posterior predictive interval contains the observed value, indicating minor deviations from the data. Since the PD model slightly overestimates draws, the ZPD model will not offer much in terms of goodness of fit. Indeed, fitting the ZPD model resulted in similar predictive values with no obvious differences.

Furthermore, we have reproduced the predictive distribution of the final table using the procedure described in Section 3.4. Results are presented in Table 1 including details from the observed final table. Minor differences were observed between expected and observed points. Predicted rankings were calculated via the expected number of points. For 60 and 80% of the competing teams, the absolute difference between predicted and observed points was found to be ≤ 2 and ≤ 3 points, respectively. Only for four teams, the corresponding difference was found between 3 and 7 points. Predictive and observed rankings are also close since, for 80% of the teams, the final observed and predicted ranking was the same or changed by only one position.

The largest differences are related to the performance of Bolton and Fulham. Bolton managed to earn 7 points additional to the ones expected (56 instead of 49) with final ranking 7 instead of the expected 12, while Fulham earned 6 points additional to the ones expected (39 instead of 33.1) getting 16th position instead of the predicted 19th avoiding by this way the relegation to the next division.

In order to examine this behaviour, we have calculated the outcome probabilities for each game of Bolton and the points expected to be earned in each game. We have traced as outliers (surprising results according to our model) all games with absolute difference between the expected and observed number of points greater than 1.95 (first criterion) or games with probability of observed outcome lower than 20% (second criterion). Six games were traced as surprising results (or outliers) as presented in Table 2.

TABLE 1 *Observed and predicted under the model points and goal differences for all the teams; in the second column (within brackets) the actual ranking is provided*

| Posterior predictive table | | | | Observed final table | | | |
|----------------------------|-------------------|------------------------|-----------------|----------------------|-------------------|-----------------|-----------------|
| Predicted (observed) | | Posterior expectations | | Observed | | Observed values | |
| Rank | Team | Points | Goal difference | Rank | Team | Points | Goal difference |
| 1 (1) | Manchester United | 86.7 | 56.0 | 1 | Manchester United | 89 | 56 |
| 2 (2) | Chelsea | 81.0 | 40.0 | 2 | Chelsea | 83 | 40 |
| 3 (3) | Arsenal | 70.5 | 28.0 | 3 | Liverpool | 68 | 30 |
| 4 (3) | Liverpool | 69.4 | 30.2 | 4 | Arsenal | 68 | 28 |
| 5 (6) | Everton | 62.5 | 16.0 | 5 | Tottenham | 60 | 3 |
| 6 (8) | Reading | 55.5 | 5.4 | 6 | Everton | 58 | 16 |
| 7 (5) | Tottenham | 54.0 | 3.2 | 7 | Bolton | 56 | -5 |
| 8 (9) | Portsmouth | 53.3 | 2.8 | 8 | Reading | 55 | 5 |
| 9 (10) | Blackburn | 51.8 | -2.1 | 9 | Portsmouth | 54 | 3 |
| 10 (11) | Aston Villa | 51.5 | 1.6 | 10 | Blackburn | 52 | -2 |
| 11 (12) | Middlesbrough | 49.0 | -4.7 | 11 | Aston Villa | 50 | 2 |
| 12 (7) | Bolton | 49.0 | -5.7 | 12 | Middlesbrough | 46 | -5 |
| 13 (13) | Newcastle | 43.8 | -8.8 | 13 | Newcastle | 43 | -9 |
| 14 (14) | Manchester City | 41.8 | -14.8 | 14 | Manchester City | 42 | -15 |
| 15 (15) | West Ham | 38.6 | -24.3 | 15 | West Ham | 41 | -24 |
| 16 (17) | Wigan | 38.1 | -21.6 | 16 | Fulham | 39 | -22 |
| 17 (18) | Sheffield United | 37.0 | -23.0 | 17 | Wigan | 38 | -22 |
| 18 (19) | Charlton | 35.7 | -25.9 | 18 | Sheffield United | 38 | -23 |
| 19 (16) | Fulham | 33.1 | -22.0 | 19 | Charlton | 34 | -26 |
| 20 (20) | Watford | 29.7 | -30.3 | 20 | Watford | 28 | -30 |

TABLE 2 *Surprising results for Bolton; games with expected absolute point difference >1.95 or probability of observed outcome <0.20*

| Home team | Away team | Final score | Goal difference (z_i) | Probabilities | | | Observed points | Expected points | Point difference |
|-------------|-----------|-------------|---------------------------|---------------|------|----------|-----------------|-----------------|------------------|
| | | | | Home win | Draw | Away win | | | |
| Bolton | Arsenal | 3-1 | 2 | 0.25 | 0.29 | 0.47 | 3 | 1.0 | 2.0 |
| Bolton | Wigan | 0-1 | -1 | 0.61 | 0.23 | 0.16 | 0 | 2.1 | -2.1 |
| Aston Villa | Bolton | 0-1 | -1 | 0.48 | 0.32 | 0.20 | 3 | 0.9 | 2.1 |
| Blackburn | Bolton | 0-1 | -1 | 0.56 | 0.21 | 0.24 | 3 | 0.9 | 2.1 |
| Portsmouth | Bolton | 0-1 | -1 | 0.53 | 0.27 | 0.20 | 3 | 0.9 | 2.1 |
| Chelsea | Bolton | 2-2 | 0 | 0.77 | 0.19 | 0.04 | 1 | 0.3 | 0.7 |

TABLE 3 *Deviations between observed and predictive measures*

| Comparison | Deviation | |
|--|-----------|-------|
| | PD | ZPD |
| 1. Relative frequencies (counts/games) | 1.06 | 1.32 |
| 2. Frequencies (counts) | 4.04 | 5.04 |
| 3. Relative frequencies of win/draw/lose | 2.20 | 2.80 |
| 4. Frequencies of win/draw/lose | 8.30 | 10.65 |
| 5. Expected points | 3.02 | 3.07 |
| 6. Expected goal difference | 0.28 | 0.40 |

Note that for Fulham, similar analysis can be implemented resulting in equivalent inference concerning surprising results.

For Bolton, we note that it earned 32 points instead of the expected 30.6 in their home field (+1.4 points) and 24 instead of the expected 18.3 in their away games (+5.7 points). Hence, Bolton performance in away games was much higher than expected. Specifically, in their home field won Arsenal where the expected number of points was equal to 1 (and probability of loosing equal to 47%) but they lost by Wigan in which game the expected number of points was equal to 2.1 and the probability of loosing only 16% (probability of winning this game 61%). All four away games with surprising scores are in favour of Bolton. They surprisingly won Aston Villa, Blackburn and Portsmouth (with probabilities ranging from 20 to 24%), while they managed to draw with Chelsea (with probability of draw equal to 19%). The additional point difference in these four games is equal to 7 points which is the difference between the predicted and observed league tables.

The same type of analysis can be performed for all teams. We have simply focused on Bolton in order to understand why its performance diverged from the one expected by the model.

Finally, Table 3 presents some goodness-of-fit measures based on the predictive distribution. Namely, for each of the quantity appearing in the table, denoted in a general form as Q , we have calculated the deviation given by

$$\text{Deviation} = \sqrt{\frac{1}{|Q|} \sum_{i=1}^{|Q|} (E(Q_i^{\text{pred}} | \mathbf{y}) - Q_i^{\text{obs}})^2},$$

where $|Q|$ is the length of vector Q . For the calculation of the deviations of the frequencies and the relative, we used $|Q| = 13$ to consider for differences from -6 to 6 , while for the deviations of the expected points and the expected differences from the final table, we set $|Q| = K = 20$, i.e. the number of teams in the league. So Q_i^{obs} is the observed quantity and $E(Q_i^{\text{pred}}|y)$ the predicted quantity; the quantities used are reported in the first column of Table 3. The results show a satisfactory fit of the model to the actual data. In addition, the zero-inflated model does not seem to improve the fit of the model for these data since we did not observe any excess of draws. We have also considered the case of a double Poisson model, where both scores and not their difference were considered. The double Poisson models provide larger expected points and goal differences mainly because it is based on a restricted Poisson assumption.

5. Concluding remarks

In this paper, we have proposed an innovative approach for modelling football data. The proposed model has some interesting advantages over the existing ones used for the same purpose. It is based on the goal differences in each game and its main feature is that it accounts for the correlation by eliminating any additive covariance. For this reason, we avoid modelling correlation or imposing assumptions about its structure as needed in any bivariate distribution (e.g. in the bivariate Poisson model). The model has a straightforward Poisson latent variable interpretation although we do not need to make assumptions about the distributions of the actual goals scored by each team. Therefore, the proposed model is quite generic and applicable to data from a wide range of football leagues in which teams with different behaviour compete. Furthermore, its parameters have a relatively easy interpretation while the parameter estimation is easier than the corresponding one for the bivariate Poisson model. The PD model is appropriate for bets like the Asian handicap where only the difference between the two teams is considered. On the other hand, by considering only the goal differences, it discards part of the available data and information and, hence, it cannot be used for modelling the final score of a game.

We have also considered a possible extension of the model by considering a zero-inflated component. Although draws were underestimated using other Poisson-related models, for the 2006–2007 English Premier League data the draws are slightly overestimated using the proposed PD model. Therefore, no zero inflation was needed.

We are currently working on further extensions of the proposed model. A first direction for extending the proposed model is to consider other distributions defined on \mathcal{Z} (see, e.g. Ong *et al.*, 2008). Secondly, variable selection techniques may be implemented to identify variables with good predictive power and by this way construct a not only more precise but also parsimonious model. In addition, Bayesian model averaging techniques can be used to improve the predictive power. As far as the zero-inflated version, when excess of draws is observed, we may also incorporate covariates in the mixing proportion p in order to predict more precisely draws and further identify additional factors that increase the probability of a draw in a football game. The model in this current form does not allow for time-varying covariates. Creating a model with such structure is under consideration.

REFERENCES

- ABRAMOWITZ, M. & STEGUN, I. A. (1974) *Handbook of Mathematical Functions*. New York: Dover.
- DIXON, M. J. & COLES, S. G. (1997) Modelling association football scores and inefficiencies in football betting market. *Appl. Stat.*, **46**, 265–280.
- IRWIN, W. (1937) The frequency distribution of the difference between two Poisson variates following the same Poisson distribution. *J. R. Stat. Soc. Ser. A*, **100**, 415–416.

- KARLIS, D. & NTZOUFRAS, I. (2003) Analysis of sports data using bivariate Poisson models. *J. R. Stat. Soc. D Stat.*, **52**, 381–393.
- KARLIS, D. & NTZOUFRAS, I. (2006) Bayesian analysis of the differences of count data. *Stat. Med.*, **25**, 1885–1905.
- LEE, A. J. (1997) Modeling scores in the premier league: is Manchester United really the best? *Chance*, **10**, 15–19.
- MAHER, M. J. (1982) Modelling association football scores. *Stat. Neerl.*, **36**, 109–118.
- ONG, S. H., SHIMIZU, K. & NG, C. M. (2008) A class of discrete distributions arising from difference of two random variables. *Comput. Stat. Data Anal.*, **52**, 1490–1499.
- SKELLAM, J. G. (1946) The frequency distribution of the difference between two Poisson variates belonging to different populations. *J. R. Stat. Soc. Ser. A*, **109**, 296.
- SPIEGELHALTER, D., THOMAS, A., BEST, N. & LUNN, D. (2003) *WinBUGS User Manual, Version 1.4*. MRC Biostatistics Unit, Institute of Public Health and Department of Epidemiology & Public Health, Imperial College School of Medicine. Available at <http://www.mrc-bsu.cam.ac.uk/bugs>. (Accessed 16 July 2007).